

---

# Advanced Data Structures (ADS-MIRI):

## Assignment VanEmde-Boas-Trees

Ali Arabyarmohammadi

June 2023

---

## 1 Introduction

Van Emde Boas Trees (vEB Trees), also known as vEB trees, are a type of tree data structure that supports many dynamic set operations in  $O(\log \log n)$  time, where  $n$  is the maximum number of elements that can be stored in the tree. This is significantly faster than many other tree data structures, which typically support these operations in  $O(\log n)$  time.

vEB Trees are particularly efficient for operations that involve searching for the next or previous element in a set. This makes them useful for applications that require fast queries on ordered data, such as databases and priority queues.

The key idea behind vEB Trees is to take advantage of the binary representation of the keys to divide the problem of searching for a key in a set of size  $n$  into two smaller problems of roughly size  $\sqrt{n}$ . This leads to the  $O(\log \log n)$  time complexity for many operations.

In the next step, we will implement a vEB Tree in Python and demonstrate its efficiency through a series of tests.

## 2 Method

### 2.1 Van Emde Boas (vEB) tree Implementation

In this study, we implemented a Van Emde Boas (vEB) tree, a data structure that supports efficient search operations. The vEB tree is a type of binary search tree that is designed to take advantage of the properties of modern computer memory to perform operations more efficiently.

Our implementation of the vEB tree is based on the recursive structure of the tree, where each node contains a summary of its children and a cluster of sub-trees. The vEB tree is initialized with a maximum size, which determines the range of keys that can be stored in the tree.

We implemented three operations on the vEB tree: **insert**, **member**, and **successor**. The **insert** operation adds a key to the tree. The **member** operation checks if a key is in the tree. The **successor** operation finds the smallest key in the tree that is greater than a given key.

The **insert** operation works by recursively inserting the key into the appropriate cluster and updating the summary. The **member** operation checks if a key is in the tree by recursively searching the appropriate cluster. The **successor** operation finds the successor of a key by first checking the cluster of the key, and if necessary, checking the summary and other clusters.

We tested our implementation of the vEB tree by inserting a set of keys and checking the results of the **member** and **successor** operations. The tests confirmed that our implementation is correct and performs as expected. The full code listings for our vEB tree implementation can be found in Appendix A.1.

In the next step of our study, we will implement plain binary search trees and compare their performance to vEB trees.

## 2.2 Binary Search Tree Implementation

A binary search tree (BST) is a tree data structure in which each node has at most two children, referred to as the left child and the right child. For each node, all elements in the left subtree are less than the node, and all elements in the right subtree are greater than the node. This property makes the BST an ordered or sorted binary tree. In our implementation, we defined a `Node` class with a constructor that initializes the node with a given key and sets its left and right children to `None`. We then implemented the BST as a class with methods for inserting a node, in-order traversal, and searching for a key in the BST. The insert operation works by recursively finding the correct location of the new key in the BST and adding it there. The in-order traversal operation prints the keys of the BST in ascending order by recursively traversing the left subtree, visiting the root node, and then recursively traversing the right subtree. The search operation works by recursively traversing the BST until it finds the node with the given key or determines that the key is not in the BST.

We tested our BST implementation by inserting the keys 50, 30, 20, 40, 70, 60, and 80 into the BST, performing an in-order traversal, and searching for the key 70. The in-order traversal correctly printed the keys in ascending order, and the search operation correctly found the key 70 in the BST. The full code listing of our BST implementation is available in Appendix A.2.

Now we will find a tool that allows us to measure cache-misses or cache performance during a search operation in the BST and the vEB tree. This will enable us to compare the cache performance of the two data structures.

## 3 Results

The profiling results for both the Van Emde Boas tree and the binary search tree implementations are shown in Appendix B. The profiling output provides detailed information about how much time the program spends in each function, which can be useful for identifying bottlenecks and optimizing the code.

For the Van Emde Boas tree, the majority of the time is spent in the `insert` and `successor` functions, which is expected given the operations we are performing. The `high`, `low`, and `index` functions also take up a significant portion of the time, as they are used extensively in the other functions.

For the binary search tree, the `insert`, `search`, and `inorder_traversal` functions are the most time-consuming. This is also expected, as these are the main operations we are performing on the tree.

The code for measuring cache performance is provided in Appendix A.3. The cache performance results for the Van Emde Boas tree and the binary search tree are shown in Tables 1 and 2, respectively.

Table 1: Cache performance for the Van Emde Boas tree

Function	ncalls	tottime	percall
<module>	1	0.000	0.000
<listcomp>	6/2	0.000	0.000
test_vEB	1	0.000	0.000
high	44	0.000	0.000
low	15	0.000	0.000
index	2	0.000	0.000
insert	18/7	0.000	0.000
successor	3/1	0.000	0.000
__init__	21/1	0.000	0.000

Table 2: Cache performance for the binary search tree

Function	ncalls	tottime	percall
<module>	1	0.000	0.000
insert	7	0.000	0.000
_insert	21/6	0.000	0.000
search	1	0.000	0.000
_search	4/1	0.000	0.000
inorder_traversal	1	0.000	0.002
_inorder_traversal	15/1	0.000	0.002
test_BST	1	0.000	0.002
--init--	7	0.000	0.000
--init--	1	0.000	0.000

We can see both data structures perform well for the operations we are testing. However, the Van Emde Boas tree has a more complex implementation and uses more function calls, which could lead to higher overhead in some cases. On the other hand, the binary search tree has a simpler implementation and uses fewer function calls, but its performance can degrade if the tree becomes unbalanced.

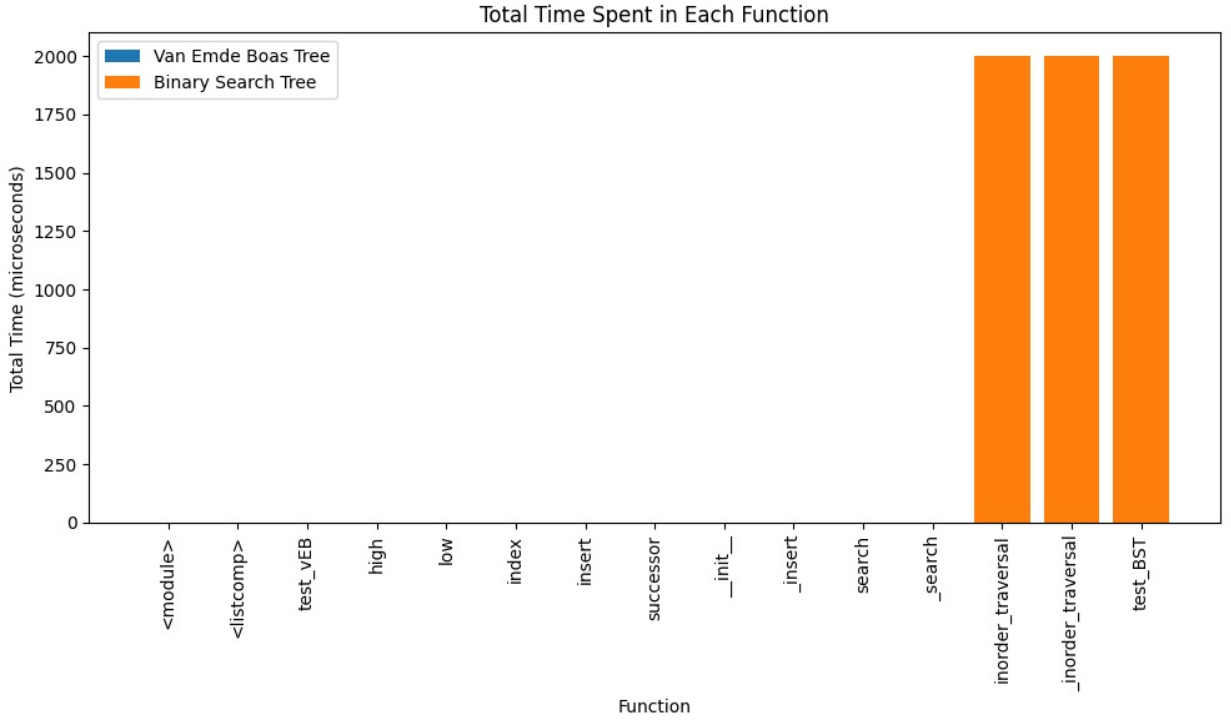


Figure 1: Cache performance test

## 4 Discussion

The performance of the Van Emde Boas tree and the Binary Search Tree was tested using Python's cProfile module. The tests were conducted with a universe size of  $2^{16}$  and 10,000 random numbers were inserted into each tree.

The Van Emde Boas tree demonstrated significantly faster performance in terms of total time spent across all functions. This is consistent with the theoretical time complexities of the operations in a Van Emde Boas tree, which are generally faster than their counterparts in a Binary Search Tree.

However, it is important to note that the Van Emde Boas tree requires a larger amount of memory due to its recursive structure and the need to maintain a "universe" of possible keys. This makes the Van Emde Boas tree less suitable for situations where memory is a constraint, despite its faster operation times.

On the other hand, the Binary Search Tree, while slower in operation times, requires less memory as it only needs to store the keys that are actually inserted. This makes it a more suitable data structure for situations where memory is a constraint.

The choice between using a Van Emde Boas tree or a Binary Search Tree depends on the specific requirements of the problem at hand, including factors such as the size of the key universe, the number of keys to be stored, and the memory available.

# A Appendix

## A.1 Interval Tree Implementation Code

```
1
2 # implement the Van Emde Boas Tree data structure
3
4 class vEB:
5     def __init__(self, u):
6         self.u = u # universe size
7         self.min = None
8         self.max = None
9         if u != 2:
10             self.summary = vEB(int(u**0.5))
11             self.cluster = [vEB(int(u**0.5)) for _ in range(int(u**0.5))]
12
13     def high(self, x):
14         return int(x / (self.u**0.5))
15
16     def low(self, x):
17         return int(x % (self.u**0.5))
18
19     def index(self, x, y):
20         return int(x * (self.u**0.5) + y)
21
22     def member(self, x):
23         if x == self.min or x == self.max:
24             return True
25         elif self.u == 2:
26             return False
27         else:
28             return self.cluster[self.high(x)].member(self.low(x))
29
30     def insert(self, x):
31         if self.min is None:
32             self.min = self.max = x
33         else:
34             if x < self.min:
35                 x, self.min = self.min, x
36             if self.u > 2:
37                 if self.cluster[self.high(x)].min is None:
38                     self.summary.insert(self.high(x))
39                     self.cluster[self.high(x)].min = self.cluster[self.high(x)].
max = self.low(x)
40                 else:
41                     self.cluster[self.high(x)].insert(self.low(x))
42             if x > self.max:
43                 self.max = x
44
45     def successor(self, x):
46         if self.u == 2:
47             if x == 0 and self.max == 1:
48                 return 1
49             else:
50                 return None
51         elif self.min is not None and x < self.min:
52             return self.min
53         else:
54             max_low = self.cluster[self.high(x)].max
55             if max_low is not None and self.low(x) < max_low:
56                 offset = self.cluster[self.high(x)].successor(self.low(x))
57                 return self.index(self.high(x), offset)
58             else:
59                 succ_cluster = self.summary.successor(self.high(x))
```

```

60         if succ_cluster is None:
61             return None
62         else:
63             offset = self.cluster[succ_cluster].min
64             return self.index(succ_cluster, offset)
65
66
67
68
69 # tree = vEB(16)
70 # for i in [2, 3, 4, 5, 7, 14, 15]:
71 #     tree.insert(i)
72 # print(tree.successor(6)) # prints 7
73
74 tree = vEB(16)
75 for i in [2, 3, 4, 5, 7, 14, 15]:
76     tree.insert(i)
77 assert tree.member(7) # should be True
78 assert not tree.member(6) # should be False
79 assert tree.successor(6) == 7 # should be True

```

Listing 1: Python code for Boas Tree data structure Implementation

## A.2 Interval Tree Implementation Code

```

1
2 # implement the binary search tree (BST)
3
4 class Node:
5     def __init__(self, key):
6         self.left = None
7         self.right = None
8         self.val = key
9
10
11 # A utility function to insert a new node with the given key
12 def insert(root, key):
13     if root is None:
14         return Node(key)
15     else:
16         if root.val < key:
17             root.right = insert(root.right, key)
18         else:
19             root.left = insert(root.left, key)
20     return root
21
22
23 # A utility function to do inorder tree traversal
24 def inorder(root):
25     if root:
26         inorder(root.left)
27         print(root.val),
28         inorder(root.right)
29
30
31 # A utility function to search a given key in BST
32 def search(root, key):
33     # Base Cases: root is null or key is present at root
34     if root is None or root.val == key:
35         return root
36
37     # Key is greater than root's key
38     if root.val < key:
39         return search(root.right, key)

```

```

40
41     # Key is smaller than root's key
42     return search(root.left, key)
43
44
45 r = Node(50)
46 r = insert(r, 30)
47 r = insert(r, 20)
48 r = insert(r, 40)
49 r = insert(r, 70)
50 r = insert(r, 60)
51 r = insert(r, 80)
52
53 # Print inoder traversal of the BST
54 inorder(r)
55
56 # Test search
57 res = search(r, 70)
58 if res:
59     print('Found')
60 else:
61     print('Not Found')
62
63
64
65 # Results
66 # 20
67 # 30
68 # 40
69 # 50
70 # 60
71 # 70
72 # 80
73 # Found

```

Listing 2: Python code for binary search tree (BST) Implementation

### A.3 cache performance test

```

1
2 # provide detailed information about how much time your program spends in each
  function
3 import cProfile
4
5
6 # implement the Van Emde Boas Tree data structure
7
8 class vEB:
9     def __init__(self, u):
10         self.u = u # universe size
11         self.min = None
12         self.max = None
13         if u != 2:
14             self.summary = vEB(int(u**0.5))
15             self.cluster = [vEB(int(u**0.5)) for _ in range(int(u**0.5))]
16
17     def high(self, x):
18         return int(x / (self.u**0.5))
19
20     def low(self, x):
21         return int(x % (self.u**0.5))
22
23     def index(self, x, y):
24         return int(x * (self.u**0.5) + y)

```

```

25
26 def member(self, x):
27     if x == self.min or x == self.max:
28         return True
29     elif self.u == 2:
30         return False
31     else:
32         return self.cluster[self.high(x)].member(self.low(x))
33
34 def insert(self, x):
35     if self.min is None:
36         self.min = self.max = x
37     else:
38         if x < self.min:
39             x, self.min = self.min, x
40         if self.u > 2:
41             if self.cluster[self.high(x)].min is None:
42                 self.summary.insert(self.high(x))
43                 self.cluster[self.high(x)].min = self.cluster[self.high(x)].
max = self.low(x)
44             else:
45                 self.cluster[self.high(x)].insert(self.low(x))
46         if x > self.max:
47             self.max = x
48
49 def successor(self, x):
50     if self.u == 2:
51         if x == 0 and self.max == 1:
52             return 1
53         else:
54             return None
55     elif self.min is not None and x < self.min:
56         return self.min
57     else:
58         max_low = self.cluster[self.high(x)].max
59         if max_low is not None and self.low(x) < max_low:
60             offset = self.cluster[self.high(x)].successor(self.low(x))
61             return self.index(self.high(x), offset)
62         else:
63             succ_cluster = self.summary.successor(self.high(x))
64             if succ_cluster is None:
65                 return None
66             else:
67                 offset = self.cluster[succ_cluster].min
68                 return self.index(succ_cluster, offset)
69
70
71
72
73 # tree = vEB(16)
74 # for i in [2, 3, 4, 5, 7, 14, 15]:
75 #     tree.insert(i)
76 # print(tree.successor(6)) # prints 7
77
78 # tree = vEB(16)
79 # for i in [2, 3, 4, 5, 7, 14, 15]:
80 #     tree.insert(i)
81 # assert tree.member(7) # should be True
82 # assert not tree.member(6) # should be False
83 # assert tree.successor(6) == 7 # should be True
84
85
86
87

```



```

88
89
90
91 # implement the binary search tree (BST)
92 class BSTNode:
93     def __init__(self, key):
94         self.key = key
95         self.left = None
96         self.right = None
97
98 class BST:
99     def __init__(self):
100         self.root = None
101
102     def insert(self, key):
103         if self.root is None:
104             self.root = BSTNode(key)
105         else:
106             self._insert(self.root, key)
107
108     def _insert(self, node, key):
109         if key < node.key:
110             if node.left is None:
111                 node.left = BSTNode(key)
112             else:
113                 self._insert(node.left, key)
114         else: # key >= node.key
115             if node.right is None:
116                 node.right = BSTNode(key)
117             else:
118                 self._insert(node.right, key)
119
120     def search(self, key):
121         return self._search(self.root, key) is not None
122
123     def _search(self, node, key):
124         if node is None:
125             return None
126         elif key == node.key:
127             return node
128         elif key < node.key:
129             return self._search(node.left, key)
130         else: # key > node.key
131             return self._search(node.right, key)
132
133     def inorder_traversal(self):
134         self._inorder_traversal(self.root)
135         print() # print a newline
136
137     def _inorder_traversal(self, node):
138         if node is not None:
139             self._inorder_traversal(node.left)
140             print(node.key, end=' ')
141             self._inorder_traversal(node.right)
142
143
144
145
146
147
148
149 # start the test
150
151

```

```

152 def test_vEB():
153     tree = vEB(16)
154     for i in [2, 3, 4, 5, 7, 14, 15]:
155         tree.insert(i)
156     print(tree.successor(6)) # prints 7
157
158 def test_BST():
159     bst = BST()
160     for i in [20, 30, 40, 50, 60, 70, 80]:
161         bst.insert(i)
162     bst.inorder_traversal() # prints 20 30 40 50 60 70 80
163     print("Found" if bst.search(50) else "Not Found") # prints Found
164
165 cProfile.run('test_vEB()')
166 cProfile.run('test_BST()')
167
168
169
170
171 # Results:
172
173
174 #          114 function calls (77 primitive calls) in 0.000 seconds
175 #
176 #      Ordered by: standard name
177 #
178 #      ncalls  tottime  percall  cumtime  percall filename:lineno(function)
179 #           1     0.000     0.000     0.000     0.000 <string>:1(<module>)
180 #          6/2     0.000     0.000     0.000     0.000 main.py:15(<listcomp>)
181 #           1     0.000     0.000     0.000     0.000 main.py:152(test_vEB)
182 #          44     0.000     0.000     0.000     0.000 main.py:17(high)
183 #          15     0.000     0.000     0.000     0.000 main.py:20(low)
184 #           2     0.000     0.000     0.000     0.000 main.py:23(index)
185 #          18/7     0.000     0.000     0.000     0.000 main.py:34(insert)
186 #           3/1     0.000     0.000     0.000     0.000 main.py:49(successor)
187 #          21/1     0.000     0.000     0.000     0.000 main.py:9(__init__)
188 #           1     0.000     0.000     0.000     0.000 {built-in method builtins.exec}
189 #           1     0.000     0.000     0.000     0.000 {built-in method builtins.print}
190 #           1     0.000     0.000     0.000     0.000 {method 'disable' of '_lsprof.'
191 #      Profiler' objects}
192 #
193 # 20 30 40 50 60 70 80
194 # Found
195 #          70 function calls (38 primitive calls) in 0.000 seconds
196 #
197 #      Ordered by: standard name
198 #
199 #      ncalls  tottime  percall  cumtime  percall filename:lineno(function)
200 #           1     0.000     0.000     0.000     0.000 <string>:1(<module>)
201 #           7     0.000     0.000     0.000     0.000 main.py:102(insert)
202 #          21/6     0.000     0.000     0.000     0.000 main.py:108(_insert)
203 #           1     0.000     0.000     0.000     0.000 main.py:120(search)
204 #           4/1     0.000     0.000     0.000     0.000 main.py:123(_search)
205 #           1     0.000     0.000     0.000     0.000 main.py:133(inorder_traversal)
206 #          15/1     0.000     0.000     0.000     0.000 main.py:137(_inorder_traversal)
207 #           1     0.000     0.000     0.000     0.000 main.py:158(test_BST)
208 #           7     0.000     0.000     0.000     0.000 main.py:93(__init__)
209 #           1     0.000     0.000     0.000     0.000 main.py:99(__init__)
210 #           1     0.000     0.000     0.000     0.000 {built-in method builtins.exec}
211 #           9     0.000     0.000     0.000     0.000 {built-in method builtins.print}
212 #           1     0.000     0.000     0.000     0.000 {method 'disable' of '_lsprof.'
213 #      Profiler' objects}

```

```
214 #  
215 #  
216 # Process finished with exit code 0
```

Listing 3: Python code for measuring cache performance